

Movie Review
Sentiment Analysis
Using Natural
Language Processing

Johnny Dryman

The IMDb logo is displayed in a bold, black, sans-serif font. The letters 'IMDb' are in a larger size than the 'b'. The logo is centered within a white rectangular box with rounded corners and a thin black border. The background of the slide features a dark gray triangle in the top left, a light gray triangle in the top right, and a yellow triangle in the bottom right, all meeting at a diagonal line.

IMDb

Goals

- Use Natural Language Processing (NLP) with machine learning models to classify IMDB user reviews as 'Positive' or 'Negative'
 - Best model – Logistic Regression with Term Frequency – Inverse Document Frequency (TF-IDF)
- Understand how our model distinguishes between 'Positive' and 'Negative'
- Explore alternative scoring system for movies based on sentiment analysis

Data

- Obtained 50,000 IMDB movie reviews
- Reviews are labeled 'positive' or 'negative,' and are considered "highly polar" reviews
 - 25,000 positive
 - 25,000 negative
- No associated movies or traditional 1-10 rankings were provided
- Model targeted 'positive' or 'negative' sentiment

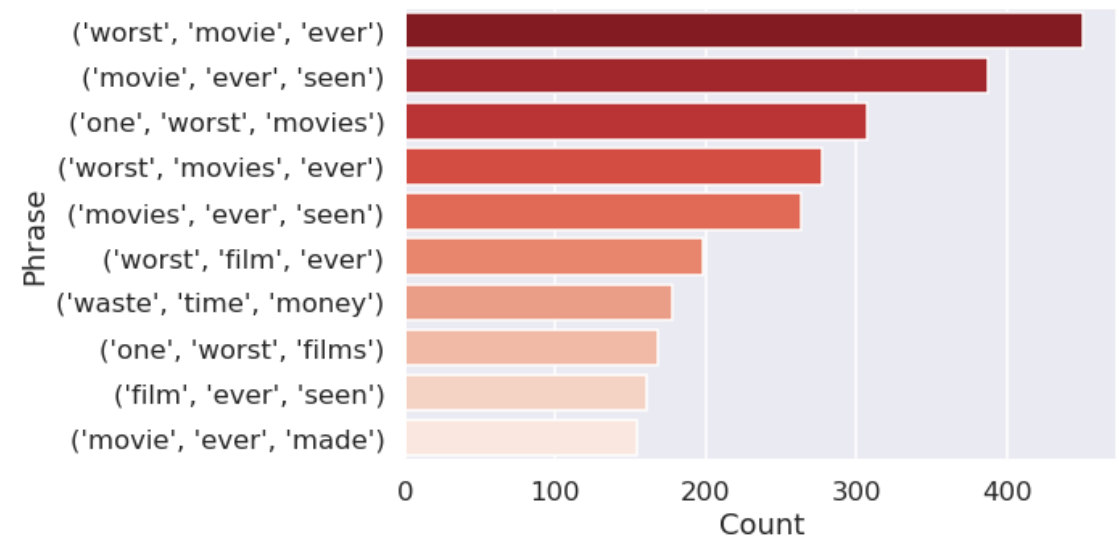
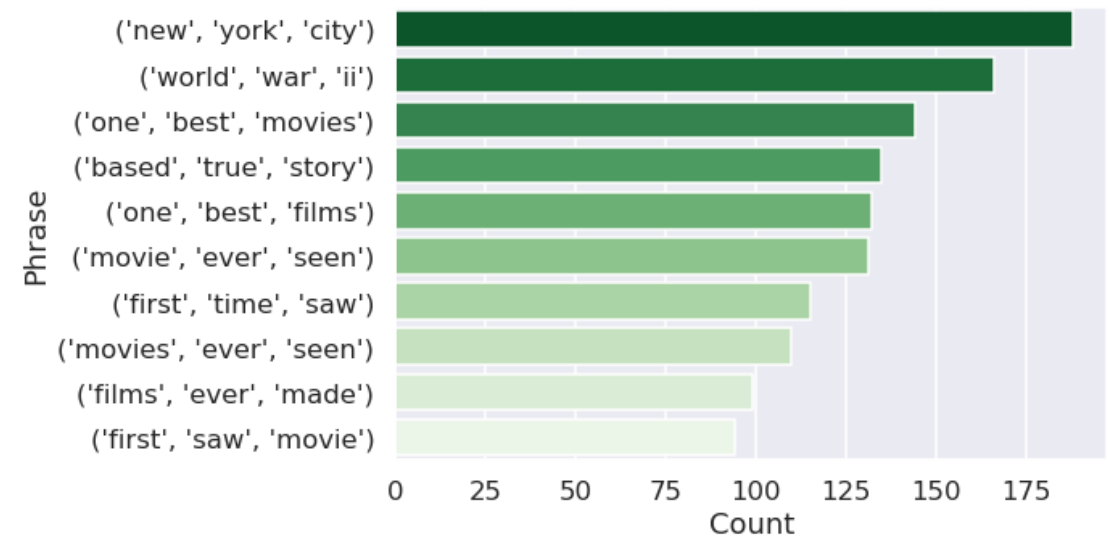
- **Trigrams**

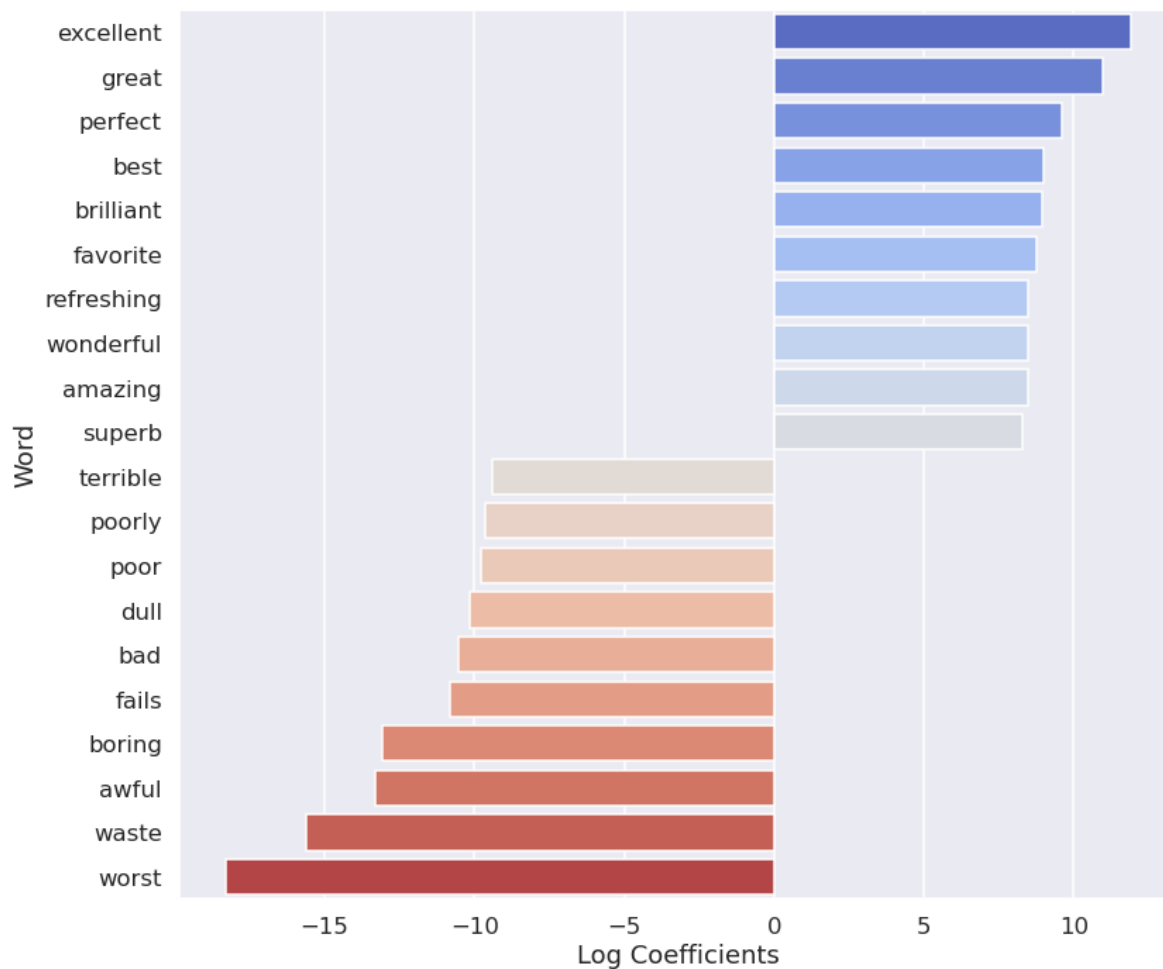
- **Positive and Negative**

- Collection of three words that appear most commonly

- Generally superlative

- “One best movies”
- “Worst film ever”





-
- **Best Model: Linear Regression**
 - Model Accuracy: **89.51%**
 - **Coefficients:**
 - **Positive and Negative**
 - Words that influenced towards or away from classifying a review as 'positive'
 - Mostly strong positive or strong negative adjectives

Conclusions

Model depended
on overtly positive
or overtly negative
adjectives

Review sentiment
can be reasonably
predicted with an
accuracy of 89.5%

Recommendations

- Use 'positive' and 'negative' sentiment model to calculate new score for all movies in database
- Compile a new list of top 250 movies
- Compare user scored top 250 with top 250 as calculated by sentiment analysis
- Depending on results, release to public

Next Steps

- Experiment with different tokenizers and stop word lists
- Reviews were added to dataset based on “highly polar” nature, quantify or qualify what this means
- Compile new dataset with buckets for multiclassification
 - Score of 1-3 / 10 = Negative
 - Score of 4-6 / 10 = Neutral
 - Score of 7-10 / 10 = Positive
- Explore alternative machine learning models
 - Support vector machines
 - K nearest neighbors
 - Deep learning / neural network models

Thanks for your time!

Please feel free to ask any questions.

